

Homework5


Ivan Ojiambo

October 26, 2017

Exe1

```
library(ggplot2)
data <- data.frame(ID      = c(1,2,3,4,5,6,7,8,9,10,11),
                   x_coord = c(3,2,3,1,4,9,5,6,6,3,3),
                   y_coord = c(9,4,3,6,1,3,6,4,2,7,5),
                   class   = c(1,1,1,1,0,0,0,0,0,0,0,"???"))

data$class <- as.factor(data$class)
#str(data)
ggplot(data, aes(x=x_coord, y=y_coord, label=ID, group = 1) ) +
  geom_point() +
  geom_text(vjust = -0.5, size = 3) +
  coord_fixed() +
  scale_x_continuous(minor_breaks = seq(0 , 10, 1), breaks = seq(0, 10, 1)) +
  scale_y_continuous(minor_breaks = seq(0 , 10, 1), breaks = seq(0, 10, 1))
```



when $K = 1$

- By looking at the graph id = 2 is the closed Neighbour to the coordinates (3, 5) therefore it belongs to class 1

when $k = 5$

- 3 => class = 1
- 4 => class = 1
- 10 => class = 0
- 7 => class = 0
- from the observation majority of the neighbouring coordinates have class = 1 , therefore class of coordinate (3, 5) = 1

when $k = 10$

- majority of the coordinated have class = 0, therefore coordinate (3, 5) = 0

Exe2


a

b. Visualise few examples of each class. Code for visualising one image is given below. Do the labels of the images correspond to what is on the image?

```
## first we need to define colors:
colors <- c('black', 'white')
cus_col <- colorRampPalette(colors=colors)

# Play around with an index of image that you want to visualise
visualise_img <- function(data,index){
  img <- array(data[index,],dim=c(28,28))
  img <- img[,28:1]
  image(img, col=cus_col(256))
}


visualise_img(mnist$x,1)
```



```
print(paste("Correct label of the first image is:", mnist$y[1]))
```

```
## [1] "Correct label of the first image is: 5"
```

```
visualise_img(mnist$x, 2)
```



```
print(paste("Correct label of the first image is:", mnist$y[2]))
```

```
## [1] "Correct label of the first image is: 0"
```

```
#yes the labels do correspond to the image
```


c

```
#c How can we visualise these labels in a bit more compact way
```

```
dataframe <- as.data.frame(mnist)
```

```
ggplot(dataframe, aes(x=y)) + geom_bar() +
```

```
  scale_x_continuous(minor_breaks = seq(0, 10, 1), breaks = seq(0, 10, 1))
```




```
##d
```

```
pixel <- data.frame('value' = mnist$x[,400],  
                    'label' = mnist$y)
```

```
ggplot(pixel, aes(x = value)) +  
  geom_bar(binwidth = 0.1) +  
  theme_bw()+  
  facet_wrap(~label)
```

```
## Warning: `geom_bar()` no longer has a `binwidth` parameter. Please use  
## `geom_histogram()` instead.
```



From the graph, + For label 0 , the 400th pixel has it has majority of

- For label 1, the 400th pixel is centered around the zero value
- FOr label 2, the black color is centered around 0 value

Exe3

a Split the data into training and testing data by filling in the ? parts in the following code:

```
# You can use set.seed() in order to reproduce stochastic results
set.seed(1111)
new_indx <- sample(c(1:nrow(mnist$x)), size = 4000, replace = FALSE)

sample_img <- mnist$x[new_indx, ]

sample_labels <- mnist$y[new_indx]

train_img    <- sample_img[1:3000,]
train_labels <- sample_labels[1:3000]
#head(train_labels)

test_img     <- sample_img[3001:4000,]
```

```
test_labels <- sample_labels[3001:4000]

str(train_img) # Make sure you have 3000 rows here

## num [1:3000, 1:784] 0 0 0 0 0 0 0 0 0 0 ...
str(test_img) # Make sure you have 1000 rows here

## num [1:1000, 1:784] 0 0 0 0 0 0 0 0 0 0 ...
```

b

```
#Next we define a distance function between two images to be euclidean distance.
#Fill in the code below so the function would calculate and return the euclidean distance between img1

dist <- function(img1, img2) {
  sqrt(sum((img1 -img2) ^2))
}

print(paste("Distance between images of class", train_labels[2], "and", train_labels[8], "is", dist(train_labels[2], train_labels[8])))

## [1] "Distance between images of class 9 and 9 is 8.82156666623062"
print(paste("Distance between images of class",train_labels[2], "and", train_labels[4], "is", dist(train_labels[2], train_labels[4])))

## [1] "Distance between images of class 9 and 6 is 10.0315864086725"
```

c

```
unknown_img <- test_img[4,]
true_label <- test_labels[4]

all_distances <- apply(train_img, 1, function(img) dist(unknown_img, img))
#head(all_distances)

#c.2. Now let's find out which image is closest to our `unknown_img`. Fill in the code.

closest_index <- which(all_distances == min(all_distances))

#c.3. Almost done, now report a label with index i in labels by filling in the code.

predicted_label <-train_labels[closest_index]


#Compare it to the true label of the first image in the test labels. Is it the same?

(predicted_label == true_label)

## [1] TRUE
print(paste("Predicted class for the first image is", predicted_label ,"and the true label is", true_label))

## [1] "Predicted class for the first image is 9 and the true label is 9"
```

```
#visualising the image
visualise_img(train_img, closest_index)
```



d

```
classify <- function(unknown_img) {
  all_distances <- apply(train_img, 1, function(img) dist(unknown_img, img))

  closest_index <- which(all_distances <= min(all_distances))
  predicted_label <- train_labels[closest_index]
  return(predicted_label)
}
#testing
print(paste("Predicted class for the first image is", classify(unknown_img), "and the true label is", t.

## [1] "Predicted class for the first image is 9 and the true label is 9"
```

e

```
classify_knn <- function(unknown_img, k = 5) {
  # This step we already know from the previous exercises
  all_distances <- apply(train_img, 1, function(img) dist(unknown_img, img))
```



```

df <- data.frame("index"= c(1:3000),
                  "dist" = all_distances)
df <- df[with(df, order(dist)), ]

# We need to get indexes of K smallest distances
# (hint: use functions order() and head())
knn = head(df$index, k)

# you can print potential predictions
#print(train_labels[knn])
# print(names(sort(table(train_labels[knn]), decreasing = TRUE)))

# Very small step is left, return the most frequently predicted label
return(names(sort(table(train_labels[knn]), decreasing = TRUE))[1])
}
classify_knn(unknown_img)

## [1] "9"
#Test this version of KNN, experiment with different `K`s

print(paste("Predicted class for the first image is", classify_knn(unknown_img, k = 5),"and the true label is", true_label))

## [1] "Predicted class for the first image is 9 and the true label is 9"

```

Exe4

a

```

test_predicted <- apply(test_img, 1, function(img) classify_knn(img))

correct_predication <- sum(as.numeric(test_predicted) - as.numeric(test_labels) ==0)
print(paste("Number of correct predication =", correct_predication ))

## [1] "Number of correct predication = 908"

#accuracy
knn_accuracy = (correct_predication/1000)
print(paste("Final accuracy of our nearest neighbor classifier is", knn_accuracy,"- not bad!"))

## [1] "Final accuracy of our nearest neighbor classifier is 0.908 - not bad!"

#for training dataset

train_predication <- vector(mode = "character", length= 3000)
for(i in 1:3000){
  unknown_img <- train_img[i,]
  #remove i th item from train_image
  testing_img <- train_img[-i,]

  #remove i th item from train_labels
  true_labels <- train_labels[-i]
}

```

```

#difference in distance
all_distances <- apply(testing_img, 1, function(img) dist(unknown_img, img))

#adding the differences in a data frame
df <- data.frame("index"= c(1:2999),
                  "dist" = all_distances)

#sort the distance in descending order
df <- df[with(df, order(dist)), ]

#return 5 rows
knn = head(df$index, 5)

train_predication[i] <- names(sort(table(train_labels[knn]), decreasing = TRUE))[1]
}

#find number of train image which are the same train labels
train_img_predicated <- sum(as.numeric(train_predication) == as.numeric(train_labels))

#calculating the accuracy proportion
knn_accuracy <- (train_img_predicated/length(train_labels))*100
print(paste("Final accuracy of our nearest neighbor classifier is", knn_accuracy,"- not bad!"))

## [1] "Final accuracy of our nearest neighbor classifier is 65.0333333333333 - not bad!"

#misclassified images


# Set an index of missclassified instance you want to examine
index = 12

miss_ind = which(test_predicted != test_labels)[index] # remember function `which` in R?

colors<-c('black','white')
cus_col<-colorRampPalette(colors=colors)

img <- array(test_img[miss_ind,],dim=c(28,28))
img <- img[,28:1]
image(img, col=cus_col(256))

```



```
print(paste("This image has a class", test_labels[miss_ind], "was incorrectly predicted as", test_predi

## [1] "This image has a class 7 was incorrectly predicted as 9"
```

b

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.4.2
```

```
## Loading required package: lattice
```

```
# We will discuss this line further in more details,
# we need it now as without it, caret tries to be very smart
# and training takes too much time...
```

```
ctrl <- trainControl(method="none", number = 1, repeats = 1)
```

```
## Warning: `repeats` has no meaning for this resampling method.
```

```
# we should use train_labels as factors, otherwise caret thinks that
# this is a regression problem
```

```
knn_predication <-function(k=5, img){
  knn_fit <- train(y = as.factor(train_labels), x = data.frame(train_img),
                  method = "knn", trControl = ctrl, tuneGrid = data.frame(k))
```

```
#Use learned nearest neighbor classifier (model) for predicting test images:
test_predicted = predict(knn_fit, data.frame(img))
```

```

    return(test_predicted)
}

test_predicted <- knn_predication(5, test_img)
print(paste("Accuracy of caret nearest neighbor classifier is", sum(test_predicted == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor classifier is 0.906"
#accuracy on training data when k = 20
test_predicted1 <- knn_predication(20, train_img)
print(paste("Accuracy of caret nearest neighbor classifier is", sum(test_predicted1 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor classifier is 0.301"
#accuracy on training data when k = 120
test_predicted1 <- knn_predication(120, train_img)
print(paste("Accuracy of caret nearest neighbor classifier is", sum(test_predicted1 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor classifier is 0.325"
#accuracy on testing data when k = 20
test_predicted1 <- knn_predication(20, test_img)
print(paste("Accuracy of caret nearest neighbor classifier is", sum(test_predicted1 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor classifier is 0.888"
#accuracy on testing data when k = 120
test_predicted1 <- knn_predication(120, test_img)
print(paste("Accuracy of caret nearest neighbor classifier is", sum(test_predicted1 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor classifier is 0.791"
#using cinfusion matrix
confusionMatrix(test_predicted, test_labels)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0    1    2    3    4    5    6    7    8    9
##      0 100    0    2    0    0    0    2    0    2    0
##      1    0   95    7    2    1    0    0    4    2    1
##      2    1    0   79    0    0    0    0    0    0    0
##      3    0    0    1   96    1    2    0    0    6    2
##      4    0    0    1    0   87    0    0    0    1    2
##      5    1    0    0    6    1   73    1    0    2    0
##      6    0    0    0    1    0    4   84    0    0    0
##      7    0    0    1    4    0    0    0  106    0    3
##      8    0    0    3    3    0    0    0    0   85    0
##      9    0    0    0    2   11    4    0    4    3  101
##
## Overall Statistics
##
##              Accuracy : 0.906
##              95% CI   : (0.8862, 0.9234)
##      No Information Rate : 0.114
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa   : 0.8954

```

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.9804    1.0000    0.8404    0.8421    0.8614    0.8795
## Specificity      0.9933    0.9812    0.9989    0.9865    0.9956    0.9880
## Pos Pred Value   0.9434    0.8482    0.9875    0.8889    0.9560    0.8690
## Neg Pred Value   0.9978    1.0000    0.9837    0.9798    0.9846    0.9891
## Prevalence       0.1020    0.0950    0.0940    0.1140    0.1010    0.0830
## Detection Rate   0.1000    0.0950    0.0790    0.0960    0.0870    0.0730
## Detection Prevalence 0.1060    0.1120    0.0800    0.1080    0.0910    0.0840
## Balanced Accuracy 0.9869    0.9906    0.9197    0.9143    0.9285    0.9338
##
##          Class: 6 Class: 7 Class: 8 Class: 9
## Sensitivity      0.9655    0.9298    0.8416    0.9266
## Specificity      0.9945    0.9910    0.9933    0.9731
## Pos Pred Value   0.9438    0.9298    0.9341    0.8080
## Neg Pred Value   0.9967    0.9910    0.9824    0.9909
## Prevalence       0.0870    0.1140    0.1010    0.1090
## Detection Rate   0.0840    0.1060    0.0850    0.1010
## Detection Prevalence 0.0890    0.1140    0.0910    0.1250
## Balanced Accuracy 0.9800    0.9604    0.9175    0.9498
```

c

```
library(caret)
##training data
rpfit <- train(y = as.factor(train_labels), x = data.frame(train_img), method = "rpart", tuneGrid = data.frame())

#accuracy of test images when using rpart method
test_predicted1 = predict(rpfit, data.frame(test_img))
print(paste("Accuracy of caret nearest neighbor using rpart", sum(test_predicted1 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor using rpart 0.623"

##confusion matrix for rpart on test images
confusionMatrix(test_predicted1, test_labels)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction    0    1    2    3    4    5    6    7    8    9
##          0  84    2    8    8    0    8    2    1    2    0
##          1    0   64    3    4    0    1    0    1    8    2
##          2    0    8   50    3    1    1    5    0    6    0
##          3    1    5    5   60    0   11    0    0    5    5
##          4    2    1    0    3   62    9    4    1    3    3
##          5    1    1    1   10    4   31    7    0    5    7
##          6    6    3   11    8    7   10   53    1   14    2
##          7    5    9    8    9    4    4    0   100    6   11
##          8    1    1    2    2    0    0    2    1   40    0
##          9    2    1    6    7   23    8   14    9   12   79
##
## Overall Statistics
```

```

##
##          Accuracy : 0.623
##          95% CI : (0.5921, 0.6531)
##    No Information Rate : 0.114
##    P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.5803
##    McNemar's Test P-Value : 2.406e-14
##
## Statistics by Class:
##
##          Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.8235   0.6737   0.5319   0.5263   0.6139   0.3735
## Specificity      0.9655   0.9790   0.9735   0.9639   0.9711   0.9607
## Pos Pred Value   0.7304   0.7711   0.6757   0.6522   0.7045   0.4627
## Neg Pred Value   0.9797   0.9662   0.9525   0.9405   0.9572   0.9443
## Prevalence       0.1020   0.0950   0.0940   0.1140   0.1010   0.0830
## Detection Rate   0.0840   0.0640   0.0500   0.0600   0.0620   0.0310
## Detection Prevalence 0.1150 0.0830 0.0740 0.0920 0.0880 0.0670
## Balanced Accuracy 0.8945   0.8263   0.7527   0.7451   0.7925   0.6671
##
##          Class: 6 Class: 7 Class: 8 Class: 9
## Sensitivity      0.6092   0.8772   0.3960   0.7248
## Specificity      0.9321   0.9368   0.9900   0.9080
## Pos Pred Value   0.4609   0.6410   0.8163   0.4907
## Neg Pred Value   0.9616   0.9834   0.9359   0.9642
## Prevalence       0.0870   0.1140   0.1010   0.1090
## Detection Rate   0.0530   0.1000   0.0400   0.0790
## Detection Prevalence 0.1150 0.1560 0.0490 0.1610
## Balanced Accuracy 0.7706   0.9070   0.6930   0.8164
##
## Accuracy of training images when using rpart method
test_predicted2 = predict(rpfit, data.frame(train_img))
print(paste("Accuracy of caret nearest neighbor using rpart", sum(test_predicted2 == test_labels)/length(test_labels)))

## [1] "Accuracy of caret nearest neighbor using rpart 0.31"
##
## confusion matrix for rpart on train images
confusionMatrix(test_predicted2, train_labels)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  0   1   2   3   4   5   6   7   8   9
##    0 239    0  27  14   2  21  12   1   7   0
##    1   1 287   8   6   1   5   0  10  18   2
##    2   1  13 151  10   2   4  17   5  12   1
##    3   2   5  24 194   9  37   5   6  17   8
##    4   1   0   4   8 197  18  18   7  14  11
##    5  10   2   6  23   7 125  16   5  15  30
##    6  10   4  36  15  27  32 184   0  48  13
##    7  10  19   9  16   5  15   1 261  12  21
##    8   4   5   6  13   0   6   4   2 124   0
##    9  14   6  29   8  30  24  17  20  27 222
##
## Overall Statistics
##

```

```
##              Accuracy : 0.6613
##              95% CI : (0.6441, 0.6783)
##      No Information Rate : 0.1137
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.6235
##  McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##              Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.81849  0.84164  0.50333  0.63192  0.70357  0.43554
## Specificity      0.96898  0.98082  0.97593  0.95804  0.97022  0.95798
## Pos Pred Value   0.73994  0.84911  0.69907  0.63192  0.70863  0.52301
## Neg Pred Value   0.98020  0.97971  0.94648  0.95804  0.96951  0.94133
## Prevalence       0.09733  0.11367  0.10000  0.10233  0.09333  0.09567
## Detection Rate   0.07967  0.09567  0.05033  0.06467  0.06567  0.04167
## Detection Prevalence 0.10767  0.11267  0.07200  0.10233  0.09267  0.07967
## Balanced Accuracy 0.89374  0.91123  0.73963  0.79498  0.83690  0.69676
##
##              Class: 6 Class: 7 Class: 8 Class: 9
## Sensitivity      0.67153  0.8233  0.42177  0.7208
## Specificity      0.93213  0.9597  0.98522  0.9350
## Pos Pred Value   0.49864  0.7073  0.75610  0.5592
## Neg Pred Value   0.96579  0.9787  0.94006  0.9670
## Prevalence       0.09133  0.1057  0.09800  0.1027
## Detection Rate   0.06133  0.0870  0.04133  0.0740
## Detection Prevalence 0.12300  0.1230  0.05467  0.1323
## Balanced Accuracy 0.80183  0.8915  0.70349  0.8279
```

Summarise the results by nicely showing all accuracies you calculated

(on both training and testing data). Comment on the results. Which model worked the best?

```
print(paste("Accuracy of test images using rpart =", sum(test_predicted1 == test_labels)/length(test_labels)))
## [1] "Accuracy of test images using rpart = 0.623"
print(paste("Accuracy of train image using rpart =", sum(test_predicted2 == test_labels)/length(test_labels)))
## [1] "Accuracy of train image using rpart = 0.31"
```

- From the result Knn model worked the best since it has a high proportions compared to rpart model